

Exceptional service in the national interest

AUTOTESTER2

Modernized Trilinos CI Testing

Samuel E. Browne, Anderson Chauphan, Joseph R. Frye, Justin M. LaPre

Trilinos User-Developer Group Meeting 10/24/2024





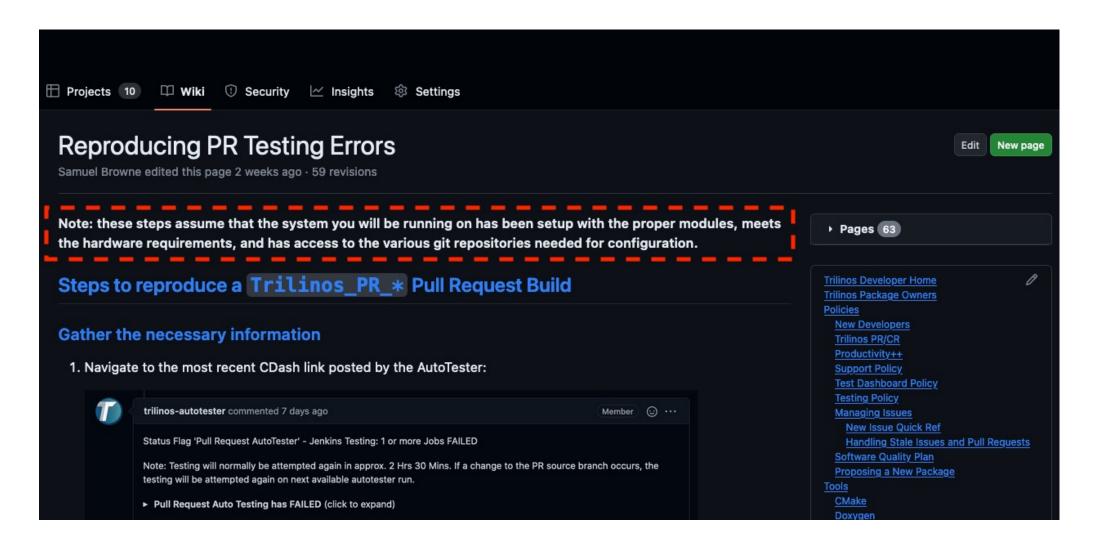






MANAGING TRILINOS BUILD ENVIRONMENTS

- Currently use a home-grown system called 'GenConfig'
- Paired with third-party library modules that are maintained on our internal systems
- How do individual developers replicate pull request builds/tests?



EXTERNALLY-UNAVAILABLE REQUIREMENTS

- 5/7 GenConfig-related repositories
- TPLs on hardware
- Hardware itself (thought we have no control over this aspect)

How can we provide the configuration tool and a software environment (TPLs) that work together to external partners?



MAKING GENCONFIG AVAILABLE TO THE COMMUNITY

- Continuing to use GenConfig (and related tooling) will require open-sourcing to make available to the broader community
- Open-source process is progressing, but is slow

GenConfig	LoadEnv	KeywordParser	SetEnvironment	DetermineSystem	ConfigParserEnhanced	SetProgramOptions

Awaiting Acceptance

Ready for GitHub

On GitHub



MAKING TPLS AVAILABLE TO THE COMMUNITY

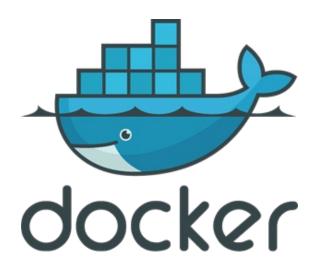
- Currently a team (SEMS) within Sandia deploys third-party libraries, compilers, and MPIs to select systems that are used for automated testing
- Un-releasable to external partners for technical reasons
 - However, SEMS moving towards delivery of TPLs with containers
- Also unavailable to internal systems outside the scope of the support agreement



CONTAINERS AS A MECHANISM FOR DISTRIBUTING TPLS

- Containers handily solve the third-party software problem
- There are limitations of reproducing novel software environments (e.g. DOE ATS systems), but these environments are not currently in pull request testing, and are outside the scope of this effort
- Containers greatly simplify the act of setting up build environments
 - Complexity is still there, but is largely handled within the Dockerfile that describes how to build the container image
 - Complexity is largely removed from user workflow





HOW TO RUN A CONTAINER

```
# Pull the image that you want to use from the registry docker pull your-registry.yourdomain.com/yourimage
```

```
# Run the image
# Remove container once it exits
# Run interactively and attach tty
# Run bash as the container entrypoint
docker run --rm -it --entrypoint bash yourimage
```



EXAMPLE

```
sebrown@triloamd01:~ $ podman run --rm -it registry-ex.sandia.gov/trilinos-project/trilinos-containers/experimental/ubi8-gcc-10.3.0
-openmpi-4.1.6:20240819
Loading gcc/10.3.0-gcc-8.5.0-ikdggsq
 Loading requirement: zlib-ng/2.1.4-gcc-8.5.0-4mix3jq zstd/1.5.5-gcc-8.5.0-4okppqr binutils/2.41-gcc-8.5.0-xt4vsa7
    gmp/6.2.1-gcc-8.5.0-w7wsbbi mpfr/4.2.0-gcc-8.5.0-3d45ev6 mpc/1.3.1-gcc-8.5.0-jdpkpms
[root@ec8674252690 /]# module list
Currently Loaded Modulefiles:
1) ccache/4.8.2
                                    11) openmpi/4.1.6
                                                                21) superlu/5.3.0
                                                                                         31) emacs/29.1
2) valgrind/3.20.0
                                    12) cmake/3.27.7
                                                                 22) superlu-dist/8.1.2 32) gh/2.32.1
                                    13) ninja/1.11.1
                                                                 23) zlib/1.3
3) qdb/13.1
4) zlib-ng/2.1.4-qcc-8.5.0-4mix3jg
                                    14) boost/1.83.0
                                                                 24) matio/1.5.17
5) zstd/1.5.5-qcc-8.5.0-4okppgr
                                     15) cgns/4.4.0
                                                                 25) libx11/1.8.4
6) binutils/2.41-gcc-8.5.0-xt4vsa7 16) hdf5/1.14.3
                                                                 26) binder/1.3.0
                                                                 27) py-mpi4py/3.1.4
7) gmp/6.2.1-gcc-8.5.0-w7wsbbi
                                    17) metis/5.1.0
8) mpfr/4.2.0-qcc-8.5.0-3d45ev6
                                    18) netcdf-c/4.9.2
                                                                 28) py-numpy/1.26.1
9) mpc/1.3.1-gcc-8.5.0-jdpkpms
                                    19) parallel-netcdf/1.12.3 29) py-pybind11/2.11.1
10) gcc/10.3.0-gcc-8.5.0-ikdggsq
                                     20) parmetis/4.0.3
                                                                 30) openblas/0.3.24
[root@ec8674252690 /]# type cmake
cmake is /home/runner/spack/opt/spack/linux-rhel8-x86_64/gcc-10.3.0/cmake-3.27.7-zxyov77bfwd7e6r5ynkwypcwqljlymw6/bin/cmake
[root@ec8674252690 /]# type ncdump
ncdump is /home/runner/spack/opt/spack/linux-rhel8-x86_64/gcc-10.3.0/netcdf-c-4.9.2-jj4cau2kiec6kcmgsd4xycyogkk4laag/bin/ncdump
[root@ec8674252690 /]#
```

Note that all TPLs are "Just There", with no module load, source, etc. All (ish) of the complexity is baked into the container recipe itself.

Can now clone Trilinos, or any other code you wish to develop.

HOW TO MOUNT YOUR LOCAL CODE INTO A CONTAINER

```
docker run --rm -it --entrypoint bash \
--mount type=bind,src=/path/on/your/machine,dst=/path/in/container \
yourimage
```

Allows you to get data in/out of container through the mounted directory Depends on host filesystem (e.g. can have some issues when mounting a Windows directory into a Linux container)

Extension: It is possible to point VSCode at a container image and have it boot said image, mount your code project for you, and then place your terminal in the running container.

CONTAINERS HELP ENSURE CONSISTENCY

Cons

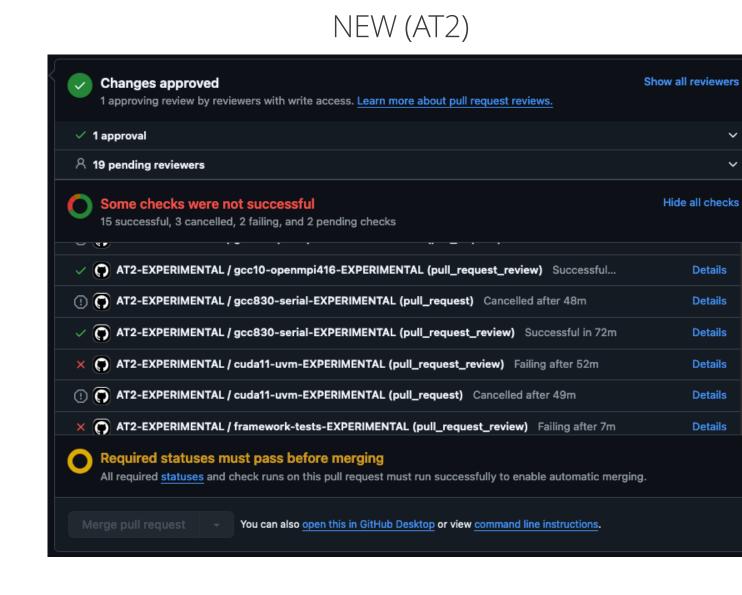
- There is overhead in learning to use containerized development environments
- Containers suffer a large performance hit for crossing CPU architectures (e.g. running an x86_64 container on an Apple Silicon MacBook)

Pros

- Near-perfect reproducibility between container runs
- Ability to easily share development environments between developers
- Anybody can create a new container on any machine with compatible architecture
- Can take container used for "validation" runs (PR testing) and run locally on developer machines



CHECKS INTERFACE



OLD (AT1)

trilinos-autotester commented on Sep 3 Member ··· Status Flag 'Pull Request AutoTester' - Jenkins Testing: 1 or more Jobs FAILED Note: Testing will normally be attempted again in approx. 2 Hrs 30 Mins. If a change to the PR source branch occurs, the testing will be attempted again on next available autotester run. ▼ Pull Request Auto Testing has FAILED (click to expand) **Build Information** Test Name: PR_gcc-openmpi-openmp • Build Num: 454 Status: FAILED Jenkins Parameters Parameter Name FORCE_CLEAN rhel8_sems-gnu-8.5.0-openmpi-4.1.6-openmp_release-debug_static_no-kokkos-GENCONFIG_BUILD_NAME $arch_no-asan_no-complex_no-fpic_mpi_no-pt_no-rdc_no-uvm_deprecated$ on_no-package-enables PR_LABELS AT: RETEST; AT: AUTOMERGE PULLREQUESTNUM 13407 PULLREQUEST_CDASH_TRACK Pull Request TEST_REPO_ALIAS **TRILINOS** TRILINOS_NODE_LABEL TRILINOS_SOURCE_REPO https://github.com/sebrowne/Trilinos TRILINOS_SOURCE_SHA TRILINOS_SRN_CONFIG TRILINOS_TARGET_BRANCH TRILINOS_TARGET_REPO https://github.com/trilinos/Trilinos TRILINOS_TARGET_SHA

(1)

. .



CDASH SNEAK PEEK

Pull Request 23 builds [view time]

			Configure		Build		Test			
Site	Build Name	Revision	Error ♥	Warn ❤	Error 💙	Warn 💙	Not Run	Fail 💙	Pass	Start Time ❤
gnu-serial-container-triloamd02	↑ PR-13528-test-rhel8_gcc-serial_release-debug_shared_no-kokkos-arch_no-asan_no-complex_no-fpic_no-mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	92e09d	0	4	0	50				12 minutes ago
cuda-container-trilogpu02	↑ PR-13407-test-rhel8_cuda-gcc-openmpi_release_static_Ampere80_no-asan_complex_no-fpic_mpi_pt_no-rdc_uvm_deprecated-on_no-package-enables	49defd	0		0	50	0	1	45	5 hours ago
gnu-openmi-container-triloamd02	↑ PR-13407-test-rhel8_gcc-openmpi_debug_shared_no-kokkos-arch_no-asan_complex_no-fpic_mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	49defd	0	4	0	50	0	0	2999	5 hours ago
gnu-serial-container-triloamd01	↑ PR-13407-test-rhel8_gcc-serial_release-debug_shared_no-kokkos-arch_no-asan_no-complex_no-fpic_no-mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	49defd	0		0	50	0	0	2211	5 hours ago
gnu-openmi-container-triloamd01	↑ PR-13527-test-rhel8_gcc-openmpi_debug_shared_no-kokkos-arch_no-asan_complex_no-fpic_mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	fbb5c6	0	3	0	50	0	0	876	3 hours ago
gnu-serial-container-triloamd01	↑ PR-13527-test-rhel8_gcc-serial_release-debug_shared_no-kokkos-arch_no-asan_no-complex_no-fpic_no-mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	fbb5c6	0		0	50	0	0	377	3 hours ago
gnu-openmi-container-triloamd02	↑ PR-13527-test-rhel8_gcc-openmpi_debug_shared_no-kokkos-arch_no-asan_complex_no-fpic_mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	fbb5c6	0	3	0	50	0	0	876	6 hours ago
gnu-serial-container-triloamd02	↑ PR-13527-test-rhel8_gcc-serial_release-debug_shared_no-kokkos-arch_no-asan_no-complex_no-fpic_no-mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	fbb5c6	0		0	50	0	0	377	6 hours ago
gnu-serial-container-triloamd01	↑ PR-13527-test-rhel8_gcc-serial_release-debug_shared_no-kokkos-arch_no-asan_no-complex_no-fpic_no-mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	fbb5c6	0	3	0	50	0	0	377	6 hours ago
gnu-openmi-container-triloamd01	A PR-13527-test-rhel8_gcc-openmpi_debug_shared_no-kokkos-arch_no-asan_complex_no-fpic_mpi_no-pt_no-rdc_no-uvm_deprecated-on_no-package-enables	fbb5c6	0		0	50	0	0	876	6 hours ago

UNAPPROVED USERS

- [] Failure Status for user dependabot[bot]: 404 Not Found
- [] Initiating User dependabot[bot] is not approved to run jobs on this machine.
- [] AT2: Please have someone from Developers trigger this workflow
- [] Job is not approved



"SPECIAL" DIRECTORIES

- [] .github directory was modified, requiring special approval...
- [] AT2: latest special approval: No special approval found
- [] AT2: Please have someone from framework review this PR and apply the AT2-SpecialApprove label

^{**} The workflow(s) will then need to be manually re-run

ACTIONS VS JENKINS-BASED CI

Cons

- Load balancing is more difficult without an "orchestration" tool (currently manually assign containers to hardware)
 - OpenShift/Kubernetes has potential to help address this

Pros

- Transparency about run stages and state
- Much more configuration-as-code under test (e.g. changing a CI configuration in the .github files is "self-testing")
- Ability to re-run only specific checks (e.g. "only the GCC check failed, it looks like a load issue, let's try re-running it")
- Jobs are queued immediately, and queued state is visible

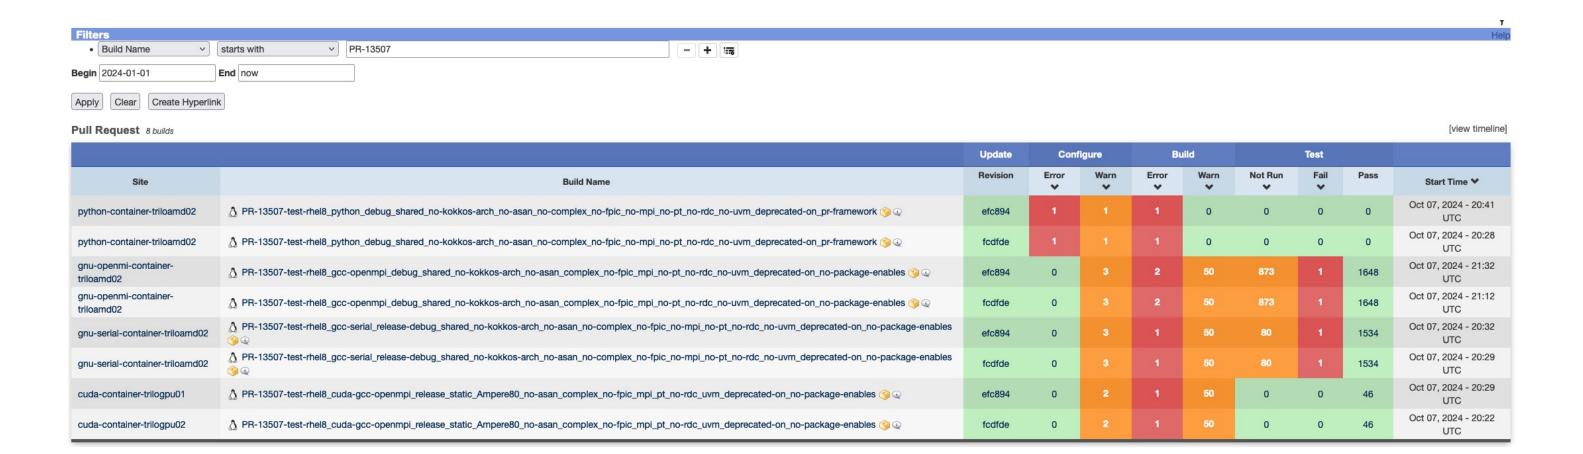


INTERACTIVE DEMO

https://github.com/trilinos/Trilinos/actions/runs/11223351354?pr=13507



FILTERED CDASH OUTPUT EXAMPLE





CONCLUSIONS

- Containers + GitHub Actions will be the CI testing driver technologies moving forwards
- Containers allow distribution of exact CI testing environments to any collaborator, internal or external
- Containers allow for testing of any containerized software stack in a similar manner (e.g. SEMS, AUE)
- GitHub Actions will allow for higher levels of transparency, hopefully fostering better developer confidence in CI
- AT2 system will allow for individual re-runs in case of system instability

Containers will be an important tool for all developers moving forwards

Thank you to the AutoTester2 team from SEMS, as well as the CSRI system admins for all of their work on enabling this system!

