



Sandia  
National  
Laboratories

# Belos Serial Dense Matrix Abstraction

Heidi Thornquist (1325) and Jennifer Loe (1465)

10/24/2024



Sandia National Laboratories is a multimission laboratory managed and operated by National Technology & Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

**SAND2024-14422PE**

# Belos Templated Linear Solvers



- Next generation linear solver libraries, written in templated C++
  - Iterative methods for solving sparse, matrix-free systems
- Provide a generic interface to a collection of algorithms for solving linear problems
- Algorithms developed with generic programming techniques
  - Algorithmic components:
    - Ease the implementation of complex algorithms
  - Operator/MultiVector interface (*and `Teuchos::ScalarTraits`*):
    - Allow the user to leverage their existing software investment
    - Multi-precision solver capability
  - Dense matrix abstract interface (**coming soon!**)
    - Allows the user to improve performance in advanced architectures
  - Design offers: Interoperability, extensibility, and reusability
- Includes block linear solvers

# Belos Solver Categories



Belos provides solvers for:

- Single RHS:  $\mathbf{Ax} = \mathbf{b}$
- Multiple RHS (available simultaneously):  $\mathbf{AX} = \mathbf{B}$
- Multiple RHS (available sequentially):  $\mathbf{Ax}_i = \mathbf{b}_i, i=1, \dots, k$
- Sequential Linear systems:  $\mathbf{A}_i \mathbf{x}_i = \mathbf{b}_i, i=1, \dots, k$
- Linear Least Squares:  $\min \|\mathbf{Ax} - \mathbf{b}\|_2$

Leverage research advances of solver community:

- Block methods: block GMRES [Vital], block CG/BICG [O'Leary]
- “Seed” solvers: hybrid GMRES [Nachtigal, et al.]
- “Recycling” solvers for sequences of linear systems: RCG, GCRO-DR [Parks, et al.]
- Restarting, orthogonalization techniques
- Multi-precision methods, communication-avoiding methods

# Belos Solvers



## Hermitian Systems ( $A = A^H$ )

- CG / Block CG
- Pseudo-Block CG (Perform single-vector algorithm simultaneously)
- RCG (Recycling Conjugate Gradients)
- PCPG (Projected CG)
- MINRES

## Non-Hermitian System ( $A \neq A^H$ )

- Block GMRES
- Pseudo-Block GMRES (Perform single-vector algorithm simultaneously)
- Block FGMRES (Variable preconditioner)
- Hybrid GMRES
- TFQMR
- GCRO-DR / Block GCRO-DR (Recycling GMRES)
- RCG (Recycling CG)

## Linear Least Squares

- LSQR

# Belos (Framework Overview)



```
 $x^{(0)}$  is an initial guess
for  $j = 1, 2, \dots$ 
  Solve  $r$  from  $Mr = b - Ax^{(0)}$ 
   $v^{(1)} = r / \|r\|_2$ 
   $s := \|r\|_2 e_1$ 
  for  $i = 1, 2, \dots, m$ 
    Solve  $w$  from  $Mw = Av^{(i)}$ 
    for  $k = 1, \dots, i$ 
       $h_{k,i} = (w, v^{(k)})$ 
       $w = w - h_{k,i} v^{(k)}$ 
    end
     $h_{i+1,i} = \|w\|_2$ 
     $v^{(i+1)} = w / h_{i+1,i}$ 
    apply  $J_1, \dots, J_{i-1}$  on  $(h_{1,i}, \dots, h_{i+1,i})$ 
    construct  $J_i$ , acting on  $i$ th and  $(i+1)$ st component
    of  $h_{\cdot,i}$ , such that  $(i+1)$ st component of  $J_i h_{\cdot,i}$  is 0
     $s := J_i s$ 
    if  $s(i+1)$  is small enough then (UPDATE( $\tilde{x}, i$ ) and quit)
  end
  UPDATE( $\tilde{x}, m$ )
end
```

GMRES Example

# Belos (Framework Overview)



Multivector  
Traits

```
 $x^{(0)}$  is an initial guess
for  $j = 1, 2, \dots$ 
  Solve  $r$  from  $Mr = b - Ax^{(0)}$ 
   $v^{(1)} = r / \|r\|_2$ 
   $s := \|r\|_2 e_1$ 
  for  $i = 1, 2, \dots, m$ 
    Solve  $w$  from  $Mw = Av^{(i)}$ 
    for  $k = 1, \dots, i$ 
       $h_{k,i} = (w, v^{(k)})$ 
       $w = w - h_{k,i} v^{(k)}$ 
    end
     $h_{i+1,i} = \|w\|_2$ 
     $v^{(i+1)} = w / h_{i+1,i}$ 
    apply  $J_1, \dots, J_{i-1}$  on  $(h_{1,i}, \dots, h_{i+1,i})$ 
    construct  $J_i$ , acting on  $i$ th and  $(i+1)$ st component
    of  $h_{\cdot,i}$ , such that  $(i+1)$ st component of  $J_i h_{\cdot,i}$  is 0
     $s := J_i s$ 
    if  $s(i+1)$  is small enough then (UPDATE( $\tilde{x}, i$ ) and quit)
  end
  UPDATE( $\tilde{x}, m$ )
end
```

GMRES Example



# Belos (Framework Overview)



## Problem Classes / Operator Traits

## Multivector Traits

```
 $x^{(0)}$  is an initial guess
for  $j = 1, 2, \dots$ 
  Solve  $r$  from  $Mr = b - Ax^{(0)}$ 
   $v^{(1)} = r / \|r\|_2$ 
   $s := \|r\|_2 e_1$ 
  for  $i = 1, 2, \dots, m$ 
    Solve  $w$  from  $Mw = Av^{(i)}$ 
    for  $k = 1, \dots, i$ 
       $h_{k,i} = (w, v^{(k)})$ 
       $w = w - h_{k,i} v^{(k)}$ 
    end
     $h_{i+1,i} = \|w\|_2$ 
     $v^{(i+1)} = w / h_{i+1,i}$ 
    apply  $J_1, \dots, J_{i-1}$  on  $(h_{1,i}, \dots, h_{i+1,i})$ 
    construct  $J_i$ , acting on  $i$ th and  $(i+1)$ st component
    of  $h_{\cdot,i}$ , such that  $(i+1)$ st component of  $J_i h_{\cdot,i}$  is 0
     $s := J_i s$ 
    if  $s(i+1)$  is small enough then (UPDATE( $\tilde{x}, i$ ) and quit)
  end
  UPDATE( $\tilde{x}, m$ )
end
```

GMRES Example

# Belos (Framework Overview)



Multivector  
Traits

Problem Classes / Operator Traits

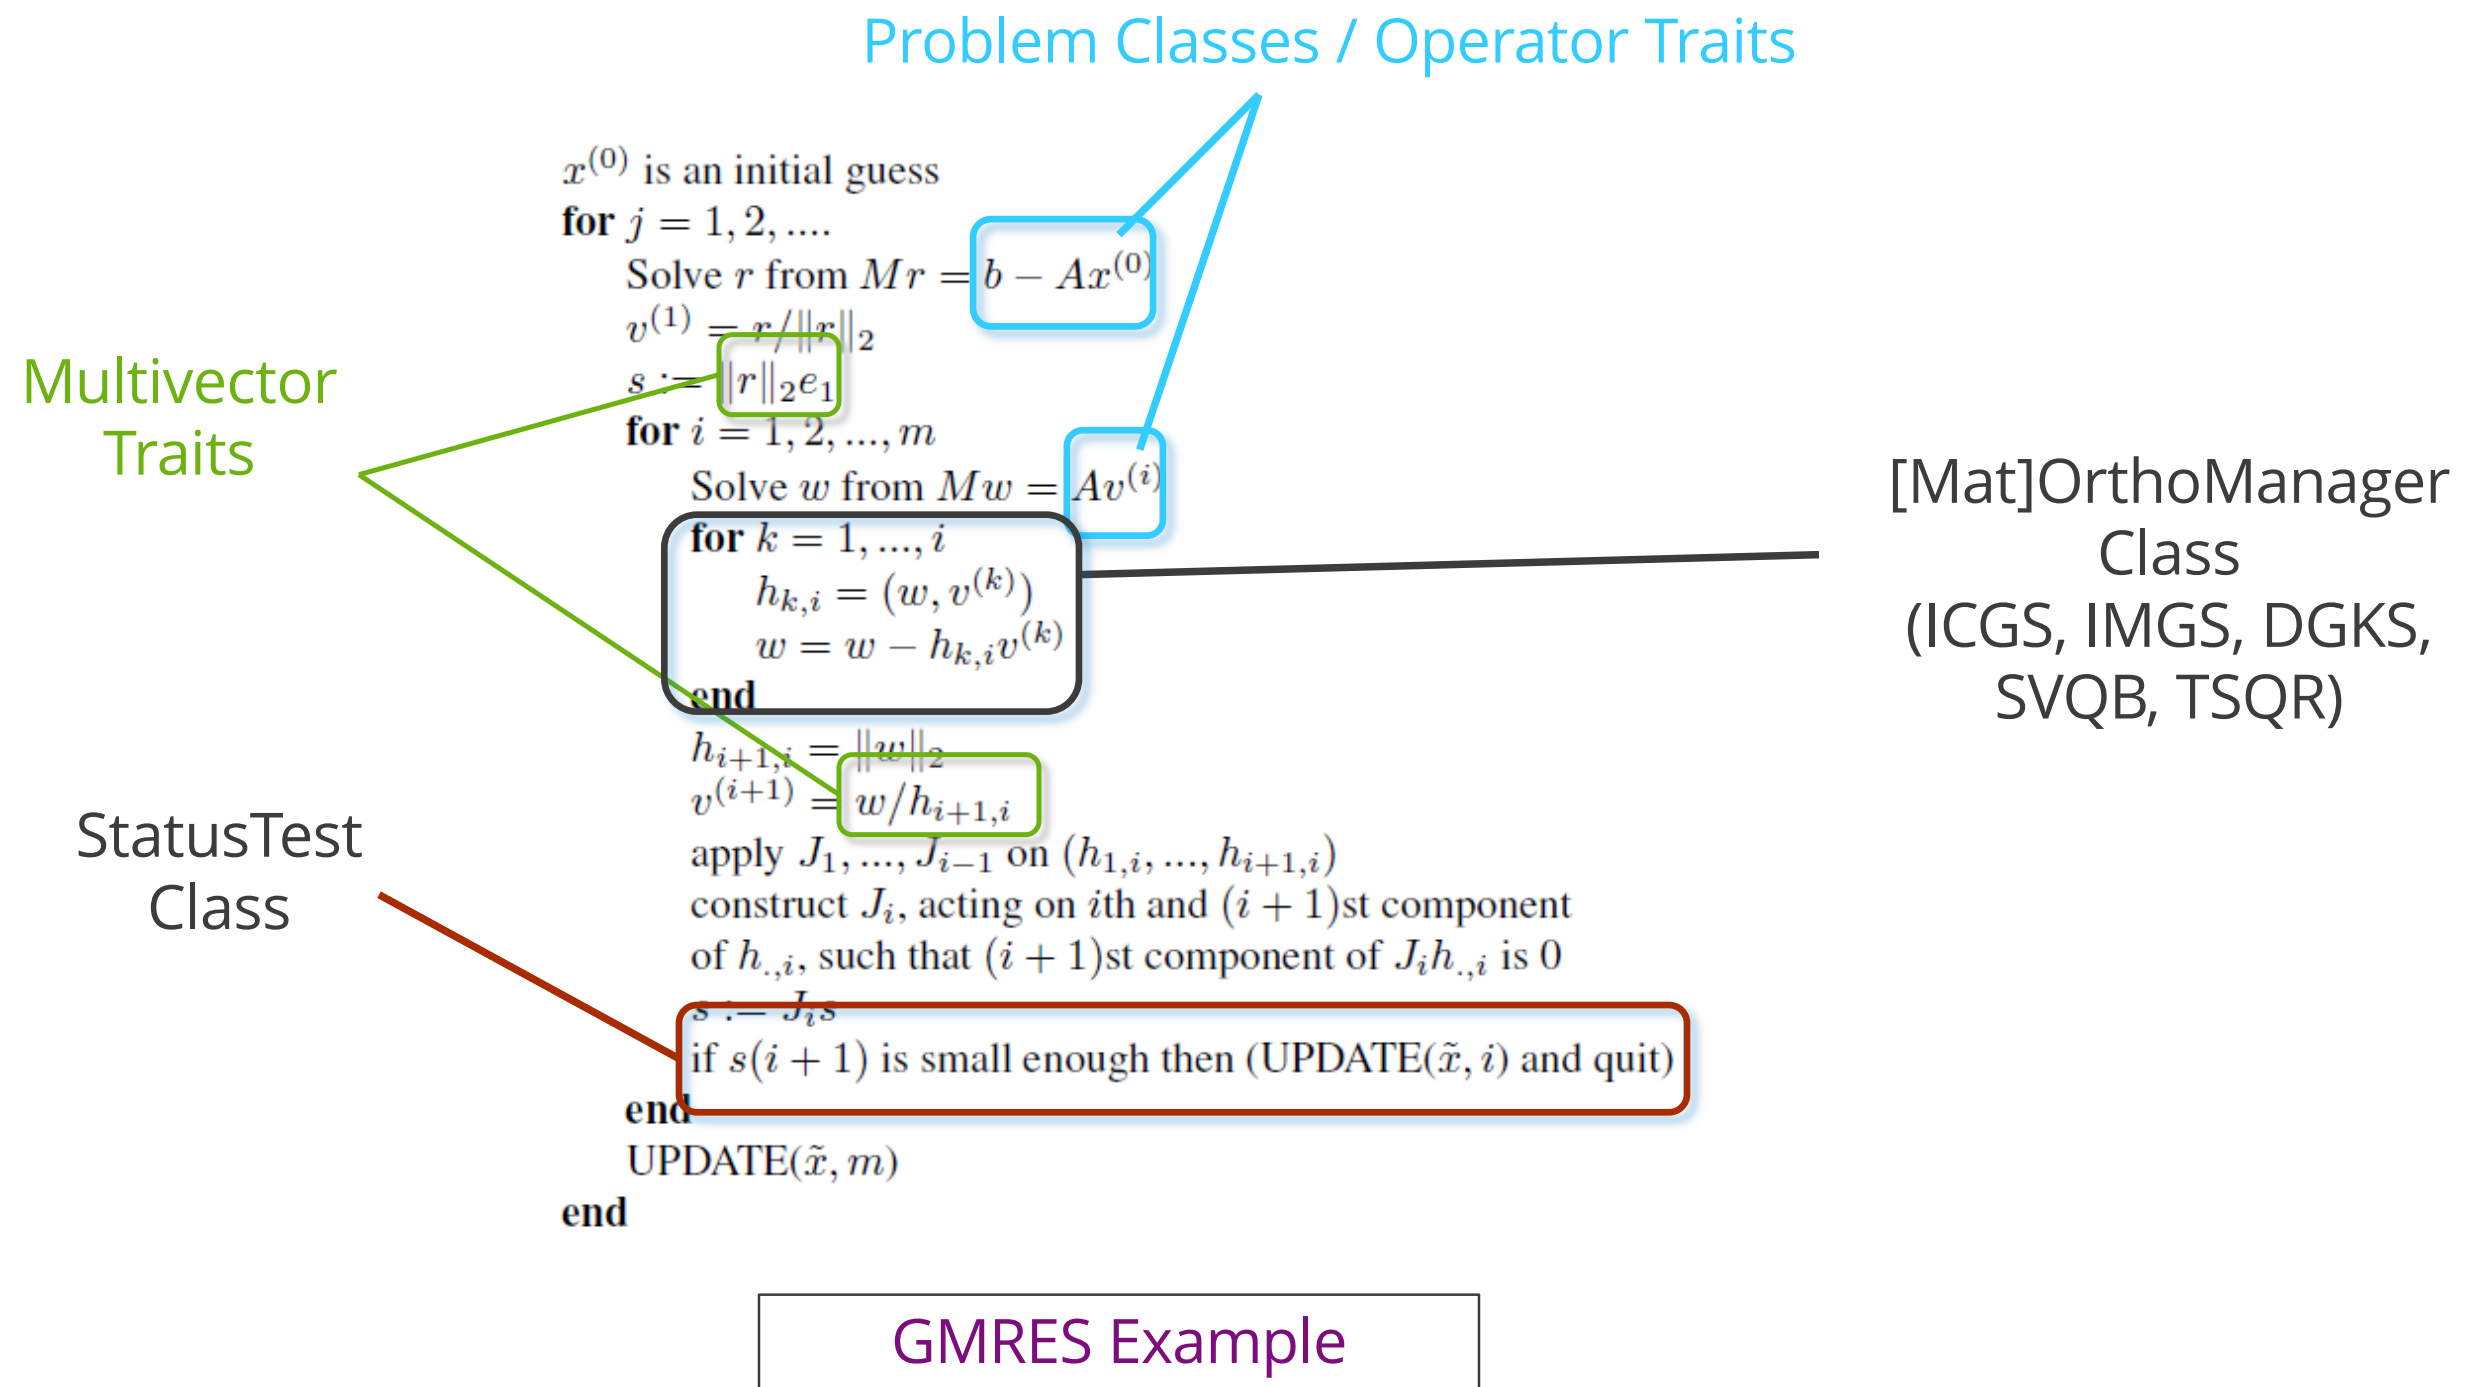
```
 $x^{(0)}$  is an initial guess
for  $j = 1, 2, \dots$ 
  Solve  $r$  from  $Mr = b - Ax^{(0)}$ 
   $v^{(1)} = r / \|r\|_2$ 
   $s := \|r\|_2 e_1$ 
  for  $i = 1, 2, \dots, m$ 
    Solve  $w$  from  $Mw = Av^{(i)}$ 
    for  $k = 1, \dots, i$ 
       $h_{k,i} = (w, v^{(k)})$ 
       $w = w - h_{k,i} v^{(k)}$ 
    end
     $h_{i+1,i} = \|w\|_2$ 
     $v^{(i+1)} = w / h_{i+1,i}$ 
    apply  $J_1, \dots, J_{i-1}$  on  $(h_{1,i}, \dots, h_{i+1,i})$ 
    construct  $J_i$ , acting on  $i$ th and  $(i+1)$ st component
    of  $h_{\cdot,i}$ , such that  $(i+1)$ st component of  $J_i h_{\cdot,i}$  is 0
     $s := J_i s$ 
    if  $s(i+1)$  is small enough then (UPDATE( $\tilde{x}, i$ ) and quit)
  end
  UPDATE( $\tilde{x}, m$ )
end
```

[Mat]OrthoManager  
Class  
(ICGS, IMGS, DGKS,  
SVQB, TSQR)

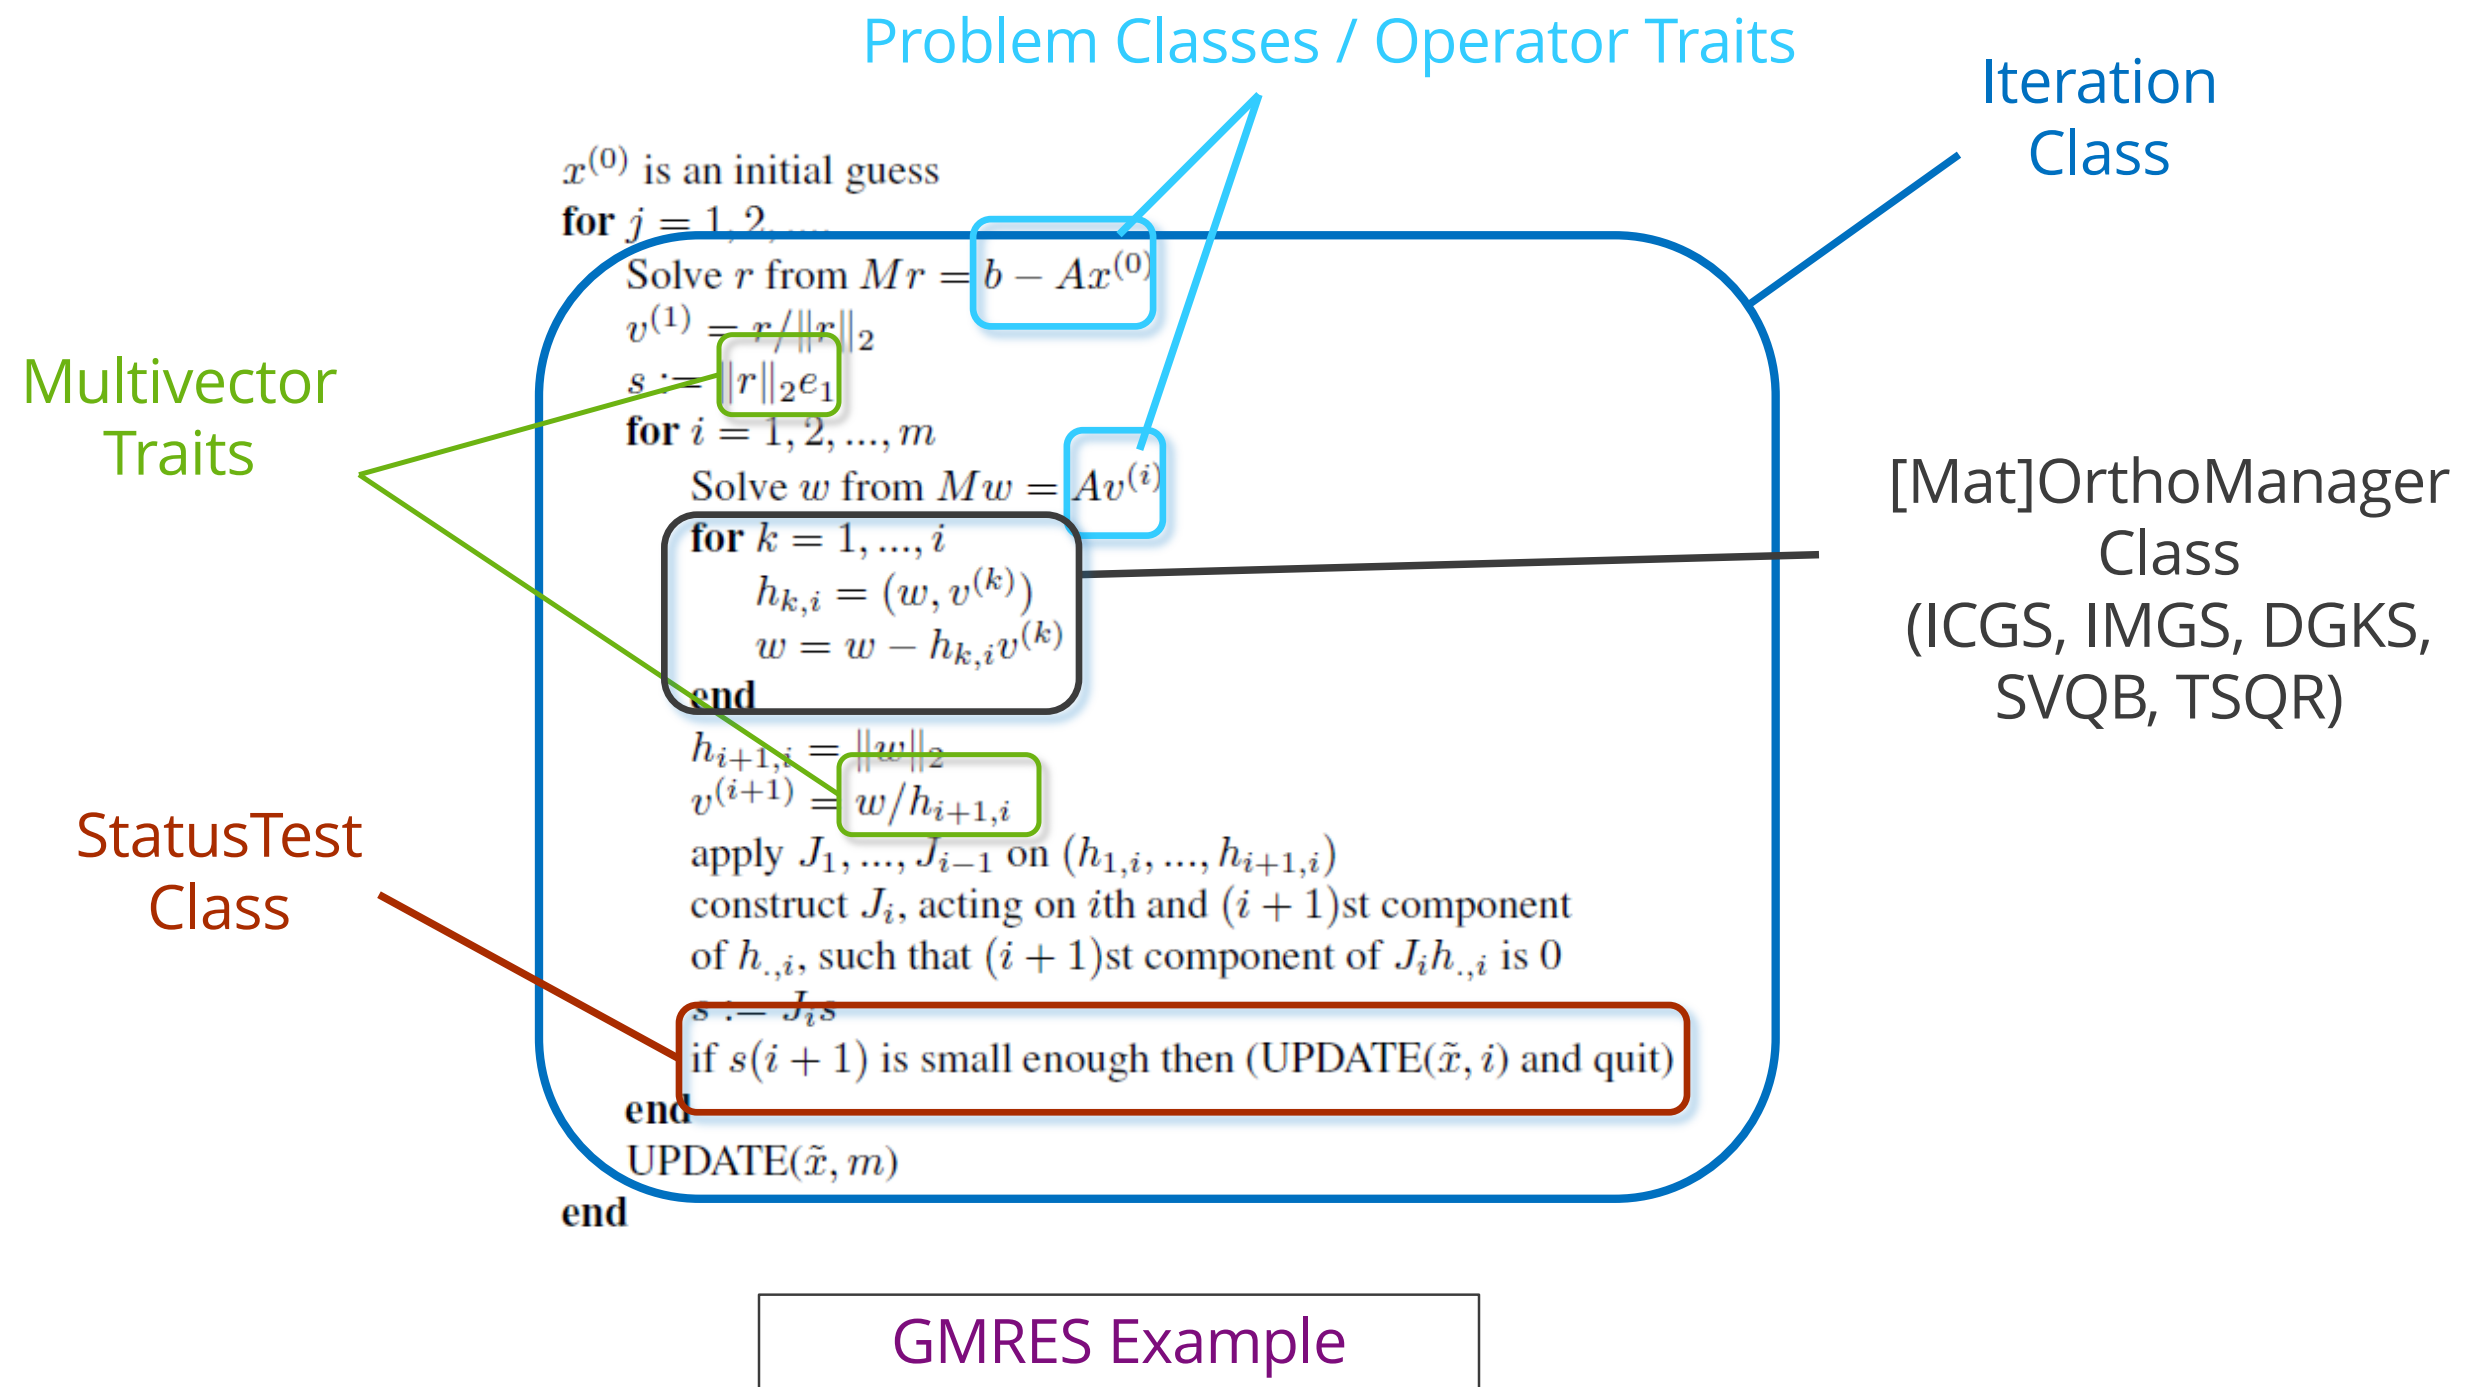
GMRES Example



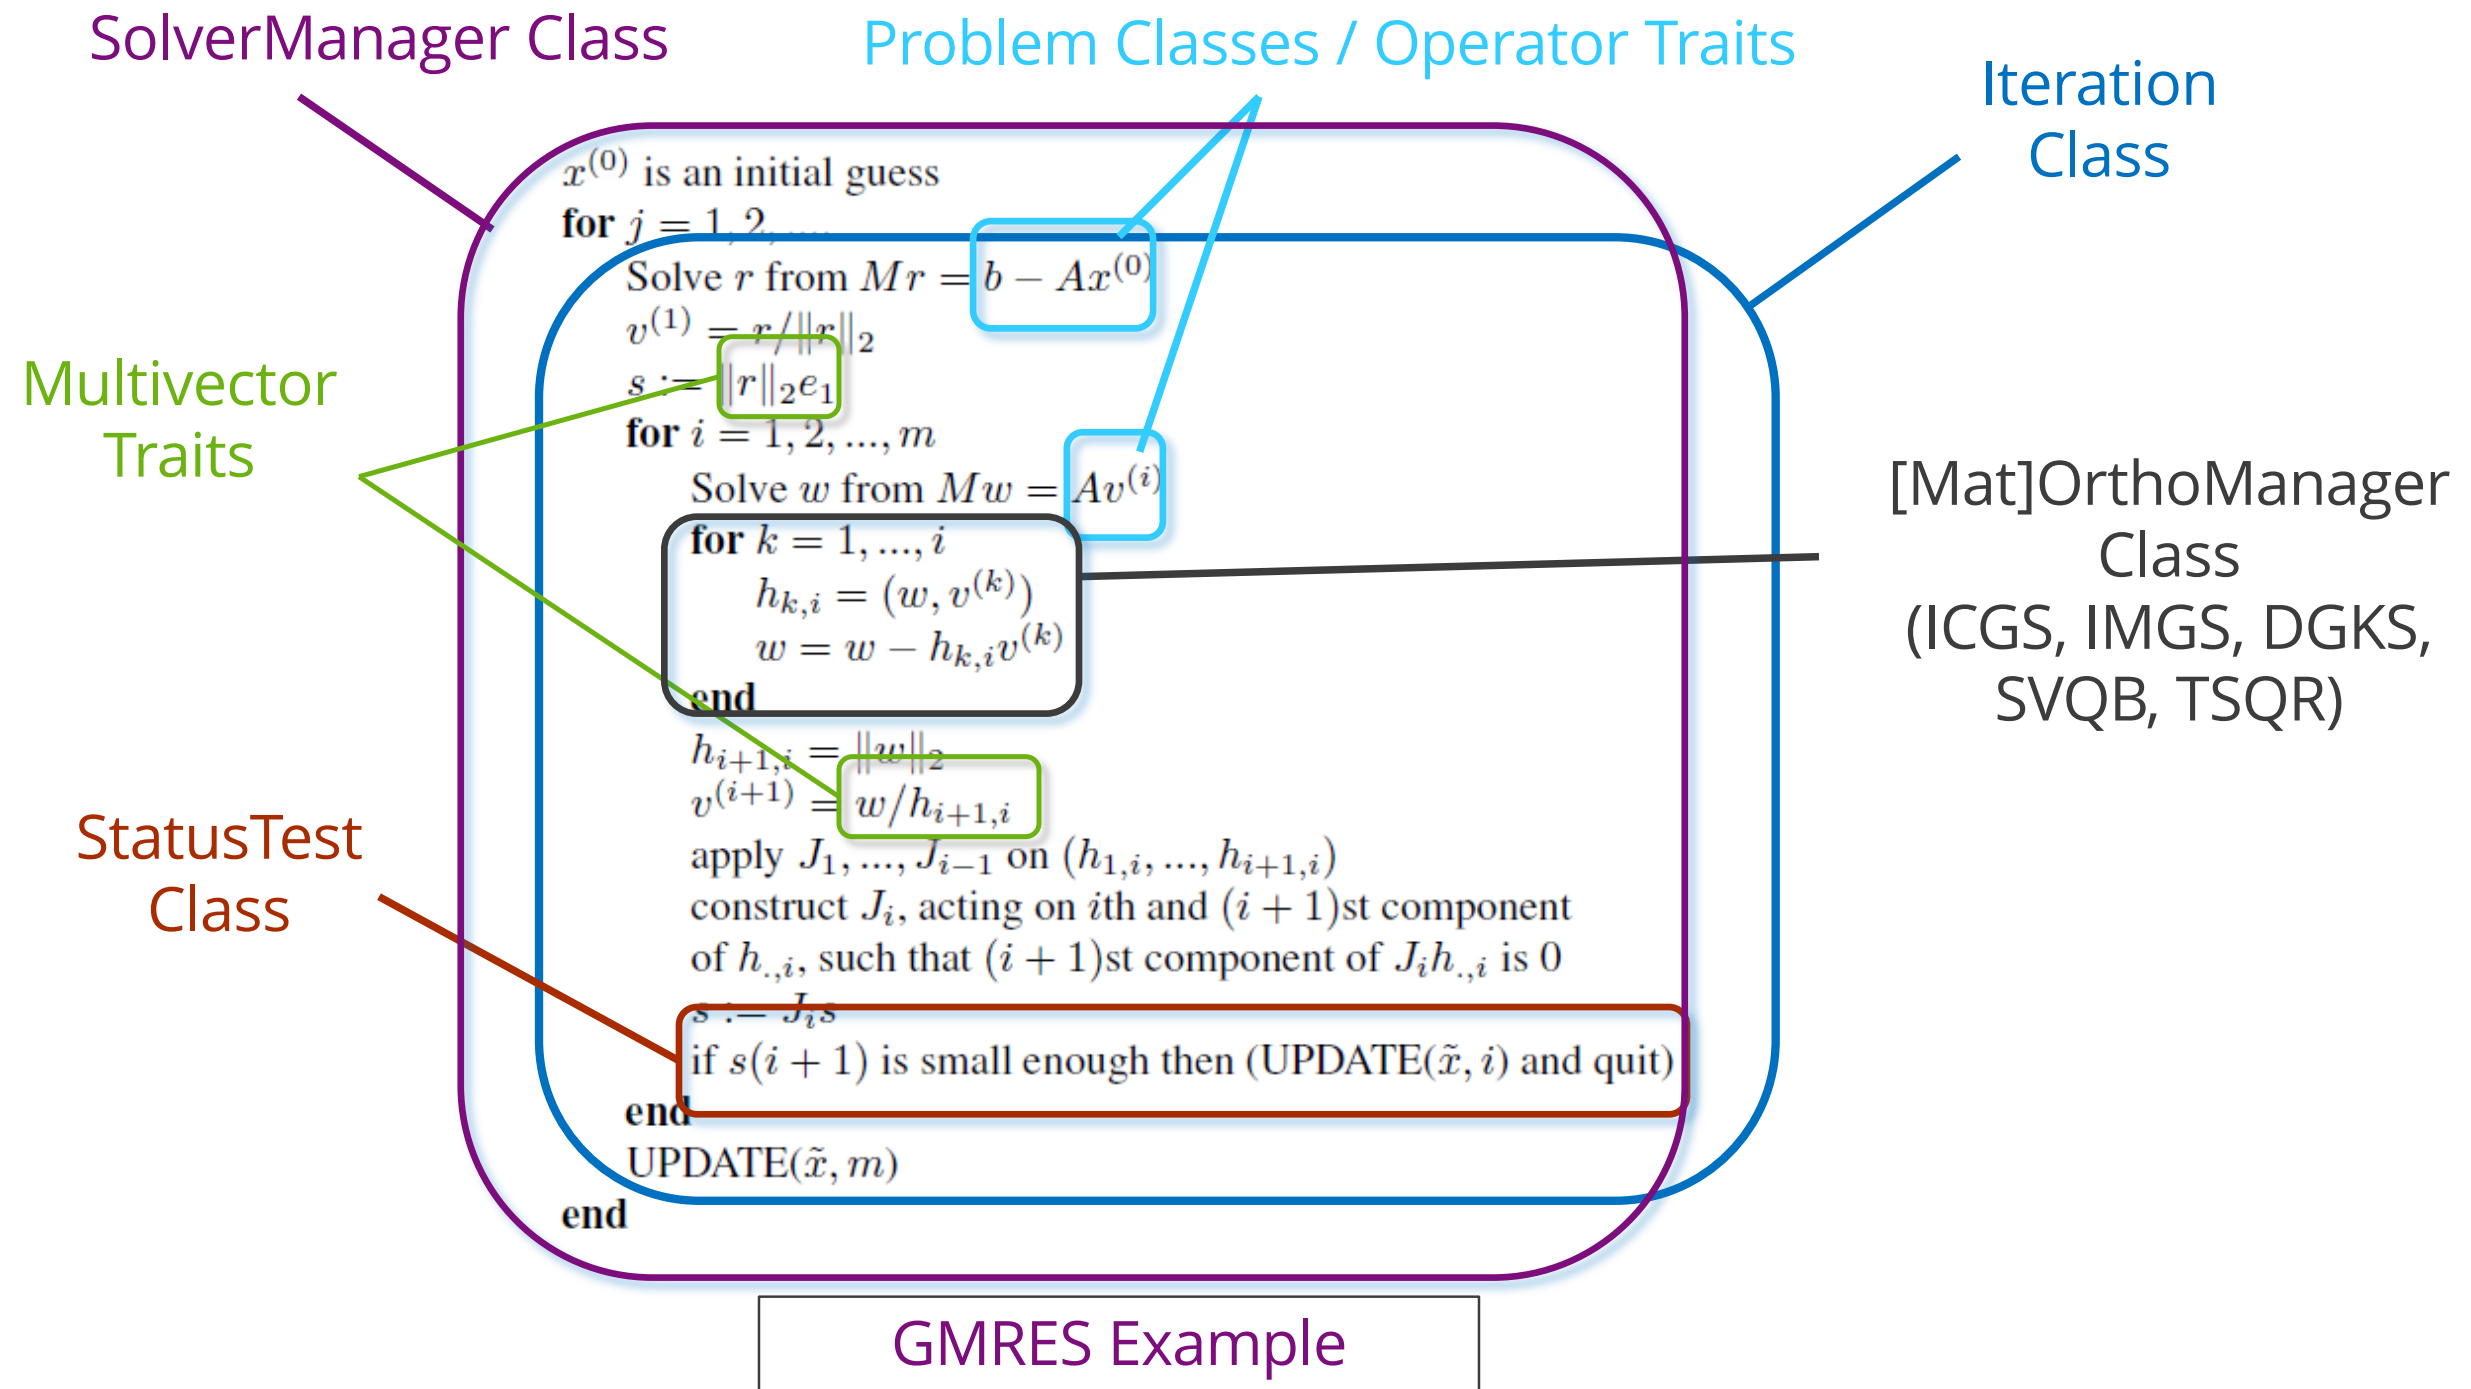
# Belos (Framework Overview)



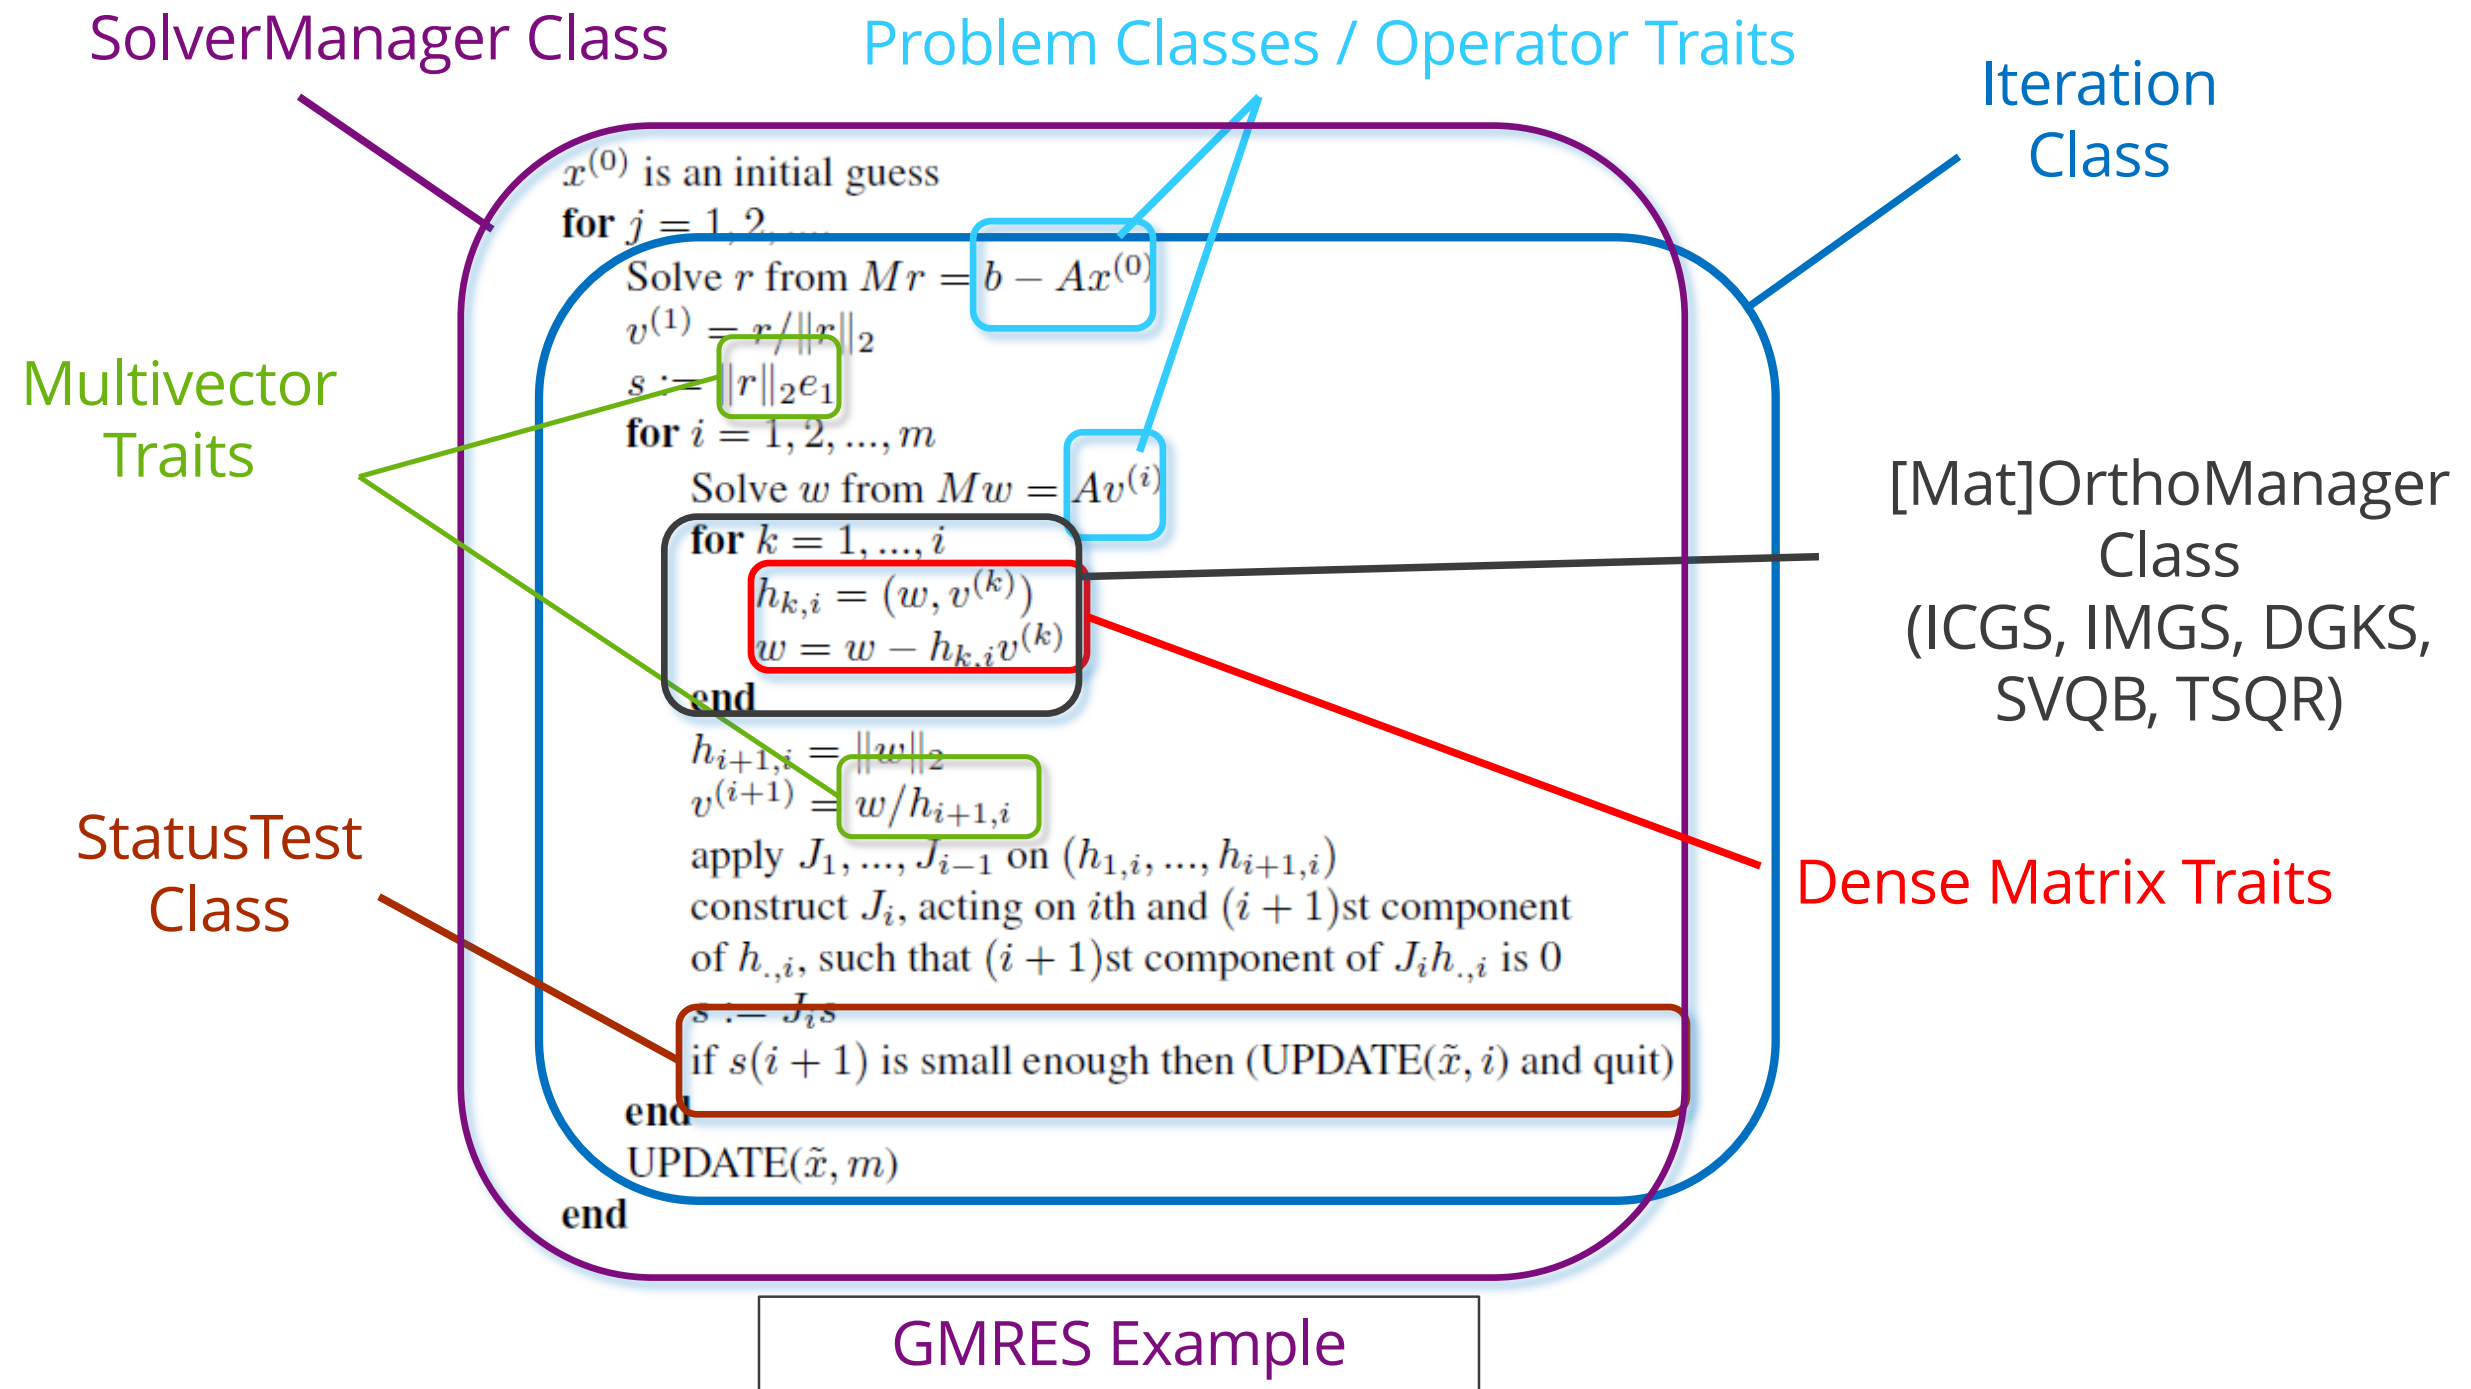
# Belos (Framework Overview)



# Belos (Framework Overview)



# Belos (Framework Overview)




# What is the Dense Matrix Traits interface?

An abstraction of the dense matrix object intended to improve iterative solver performance in advanced architectures

- Default: `Teuchos::SerialDenseMatrix<>`
- Optionally: `Kokkos::DualView<>`

So, what is the minimal interface for Belos linear solvers?



# What is the Dense Matrix Traits interface?




```
static Teuchos::RCP<DM> Create()  
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)  
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
```





# What is the Dense Matrix Traits interface?


```
static Teuchos::RCP<DM> Create()  
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)  
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)  
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
```



# What is the Dense Matrix Traits interface?



```
static Teuchos::RCP<DM> Create()
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> Subview( DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<const DM> SubviewConst( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<DM> SubviewCopy( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
```



# What is the Dense Matrix Traits interface?




```
static Teuchos::RCP<DM> Create()
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> Subview( DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<const DM> SubviewConst( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<DM> SubviewCopy( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static ScalarType* GetRawHostPtr(DM & dm )
static ScalarType const * GetConstRawHostPtr(const DM & dm )
```



# What is the Dense Matrix Traits interface?



```
static Teuchos::RCP<DM> Create()
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> Subview( DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<const DM> SubviewConst( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<DM> SubviewCopy( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static ScalarType* GetRawHostPtr(DM & dm )
static ScalarType const * GetConstRawHostPtr(const DM & dm )
static ScalarType & Value( DM& dm, const int i, const int j )
static const ScalarType & ValueConst( const DM& dm, const int i, const int j )
```



# What is the Dense Matrix Traits interface?




```
static Teuchos::RCP<DM> Create()
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> Subview( DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<const DM> SubviewConst( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<DM> SubviewCopy( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static ScalarType* GetRawHostPtr(DM & dm )
static ScalarType const * GetConstRawHostPtr(const DM & dm )
static ScalarType & Value( DM& dm, const int i, const int j )
static const ScalarType & ValueConst( const DM& dm, const int i, const int j )
static void SyncDeviceToHost(DM & dm)
static void SyncHostToDevice(DM & dm)
```



# What is the Dense Matrix Traits interface?



```
static Teuchos::RCP<DM> Create()
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> Subview( DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<const DM> SubviewConst( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<DM> SubviewCopy( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static ScalarType* GetRawHostPtr(DM & dm )
static ScalarType const * GetConstRawHostPtr(const DM & dm )
static ScalarType & Value( DM& dm, const int i, const int j )
static const ScalarType & ValueConst( const DM& dm, const int i, const int j )
static void SyncDeviceToHost(DM & dm)
static void SyncHostToDevice(DM & dm)
static int GetNumRows( const DM& dm )
static int GetNumCols( const DM& dm )
static int GetStride( const DM& dm )
```








# What is the Dense Matrix Traits interface?

```
static Teuchos::RCP<DM> Create()
static Teuchos::RCP<DM> Create( const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> CreateCopy(const DM & dm, bool transpose=false)
static void Reshape( DM& dm, const int numRows, const int numcols, bool initZero = true)
static Teuchos::RCP<DM> Subview( DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<const DM> SubviewConst( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static Teuchos::RCP<DM> SubviewCopy( const DM & source, int numRows, int numCols, int startRow=0, int startCol=0)
static ScalarType* GetRawHostPtr(DM & dm )
static ScalarType const * GetConstRawHostPtr(const DM & dm )
static ScalarType & Value( DM& dm, const int i, const int j )
static const ScalarType & ValueConst( const DM& dm, const int i, const int j )
static void SyncDeviceToHost(DM & dm)
static void SyncHostToDevice(DM & dm)
static int GetNumRows( const DM& dm )
static int GetNumCols( const DM& dm )
static int GetStride( const DM& dm )
static void Add( DM& thisDM, const DM& sourceDM)
static void PutScalar( DM& dm, ScalarType value = Teuchos::ScalarTraits<ScalarType>::zero())
static void Scale( DM& dm, ScalarType value)
static void Randomize( DM& dm)
static void Assign( DM& thisDM, const DM& sourceDM)
static typename Teuchos::ScalarTraits<ScalarType>::magnitudeType NormFrobenius( DM& dm)
static typename Teuchos::ScalarTraits<ScalarType>::magnitudeType NormOne( DM& dm)
```



# What are the impacts of this change internal to Belos?



Synchronizations are explicitly integrated into Belos solvers and support classes

- Mostly to perform some computation on the dense matrix object
- Sometimes to provide debug information (optional)

\* Uses MvDot

| Class   | # SyncDeviceToHost | Reason   |
|---|--------------------|--|
| ICGS / IMGS / DGKS                              | 1*                 | Debugging  |
| BiCGStab / TFQMR                                | 0*                 |  |
| Block CG  | 2                  | Cholesky solve (POTRF/POTRS)   |
| CG / PseudoBlock CG / PseudoBlock Stochastic CG | 2*                 | Extract scalar coefficients  |
| PCPG  | 6                  | Extract scalar coefficients (GESVD)  |
| RCG   | 48*                | Recycling subspace (GEMM/GESV/GETRS/SYGV)  |
| Minres  | 3                  | Extract scalar coefficients  |
| Block (F)GMRES / PseudoBlock GMRES              | 3                  | Update / solve LSQR (ROT/TRSM)   |
| GCRODR  | 21                 | Recycling subspace, update / solve LSQR (ROT/TRSM/ GEMM/GETRF/GETRI/GEQRF/UNGQR/GEEV/GESV/GGEVX) |
| LSQR  | 2                  | Debugging  |

# What are the impacts of this change to Belos MV Traits?



The MultiVectorTraits interface has methods with dense matrix/vector object arguments:

```
static void MvTimesMatAddMv( const ScalarType alpha, const MV& A,  
                             const DM& B,  
                             const ScalarType beta, MV& mv )
```

```
static void MvTransMv( const ScalarType alpha, const MV& A, const MV& mv, DM& B)
```

# What are the impacts of this change to Belos MV Traits?



The MultiVectorTraits interface has methods with dense matrix/vector object arguments:

```
static void MvTimesMatAddMv( const ScalarType alpha, const MV& A,  
                             const DM& B,  
                             const ScalarType beta, MV& mv )
```

```
static void MvTransMv( const ScalarType alpha, const MV& A, const MV& mv, DM& B)
```

What about this one?

```
static void MvDot ( const MV& mv, const MV& A, std::vector<ScalarType> &b)
```

Will a synchronization be hidden within this call?

```
Teuchos::ArrayView<ScalarType> av (b);
```

```
mv.dot (A, av (0, numVecs));
```

# What are the impacts of this change external to Belos?



All of the classes that directly or indirectly use a dense matrix object will have an additional template argument

```
template<class ScalarType, class MV, class OP, class DM = Teuchos::SerialDenseMatrix<int,ScalarType>>
```

# What are the impacts of this change external to Belos?



All of the classes that directly or indirectly use a dense matrix object will have an additional template argument

```
template<class ScalarType, class MV, class OP, class DM = Teuchos::SerialDenseMatrix<int,ScalarType>>
```

Since the default template argument is `Teuchos::SerialDenseMatrix<>`, this change should not be substantially impactful ... initially



# What are the impacts of this change external to Belos?



All of the classes that directly or indirectly use a dense matrix object will have an additional template argument

```
template<class ScalarType, class MV, class OP, class DM = Teuchos::SerialDenseMatrix<int,ScalarType>>
```

Since the default template argument is `Teuchos::SerialDenseMatrix<>`, this change should not be substantially impactful ... initially

Partial template specializations of Belos classes may need to be modified

- Stokhos, MueLu ☒

# What are the impacts of this change external to Belos?



All of the classes that directly or indirectly use a dense matrix object will have an additional template argument

```
template<class ScalarType, class MV, class OP, class DM = Teuchos::SerialDenseMatrix<int,ScalarType>>
```

Since the default template argument is `Teuchos::SerialDenseMatrix<>`, this change should not be substantially impactful ... initially

Partial template specializations of Belos classes may need to be modified

- Stokhos, MueLu ☒

Direct use of Belos package appears easiest way to employ this new capability

- Thyra interfaces used by Stratimikos may not be compatible with `Kokkos::DualView`?

# What is next for the Dense Matrix Abstraction?



Development and integration of this change has been completed on a fork of Trilinos

- Detailed testing of abstract interface and implementations ☒
- Stokhos, MueLu ☒

# What is next for the Dense Matrix Abstraction?



Development and integration of this change has been completed on a fork of Trilinos

- Detailed testing of abstract interface and implementations ☒
- Stokhos, MueLu ☒

Finalizing the refactor

- Fix remaining testing issues
- Obtain initial GPU performance results

# What is next for the Dense Matrix Abstraction?



Development and integration of this change has been completed on a fork of Trilinos

- Detailed testing of abstract interface and implementations ☒
- Stokhos, MueLu ☒

Finalizing the refactor

- Fix remaining testing issues
- Obtain initial GPU performance results

Future directions

- Replace synchronizations + BLAS/LAPACK methods with KokkosKernels methods

# What is next for the Dense Matrix Abstraction?



Development and integration of this change has been completed on a fork of Trilinos

- Detailed testing of abstract interface and implementations ☒
- Stokhos, MueLu ☒

Finalizing the refactor

- Fix remaining testing issues
- Obtain initial GPU performance results

Future directions

- Replace synchronizations + BLAS/LAPACK methods with KokkosKernels methods

Comments / Questions ?